
CIS 471/571 (Fall 2020): Introduction to Artificial Intelligence

Assignment Project Exam Help

Lecture 7: Expectimax, Utilities

<https://powcoder.com>

Add WeChat powcoder

Thanh H. Nguyen

Source: <http://ai.berkeley.edu/home.html>



Reminders

- Project 2: Multi-agent Search

- Deadline: Oct 27th, 2020

Assignment Project Exam Help

- Homework 2: CSPs and Games <https://powcoder.com>

- Deadline: Oct 24th, 2020 [Add WeChat powcoder](#)

Today

- Expectimax Search

Assignment Project Exam Help

- Utilities

<https://powcoder.com>


Add WeChat powcoder

Uncertain Outcomes


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



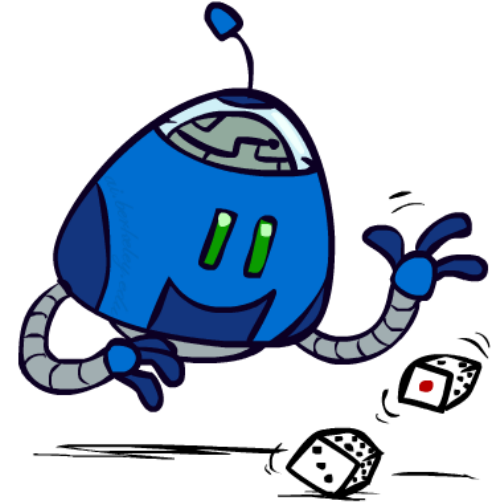

Worst-Case vs. Average Case



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Idea: Uncertain outcomes controlled by chance, not an adversary!




Expectimax Search

- Why wouldn't we know what the result of an action will be?
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the ghosts respond randomly
 - Actions can fail: when moving a robot, wheels might slip

- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes

- Expectimax search:** compute the average score under optimal play

- Max nodes as in minimax search
 - Chance nodes are like min nodes but the outcome is uncertain
 - Calculate their **expected utilities**
 - I.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as **Markov Decision Processes**



Expectimax Pseudocode

```
def value(state):
```

```
    if the state is a terminal state: return the state's utility
```

```
    if the next agent is MAX: return max_value(state)
```

```
    if the next agent is EXP: return exp_value(state)
```

<https://powcoder.com>

Add WeChat powcoder

```
def max_value(state):
```

```
    initialize v =  $-\infty$ 
```

```
    for each successor of state:
```

```
        v = max(v, value(successor))
```

```
    return v
```

```
def exp_value(state):
```

```
    initialize v = 0
```

```
    for each successor of state:
```

```
        p = probability(successor)
```

```
        v += p * value(successor)
```

```
    return v
```


Expectimax Pseudocode

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```

Assignment Project Exam Help

<https://powcoder.com>

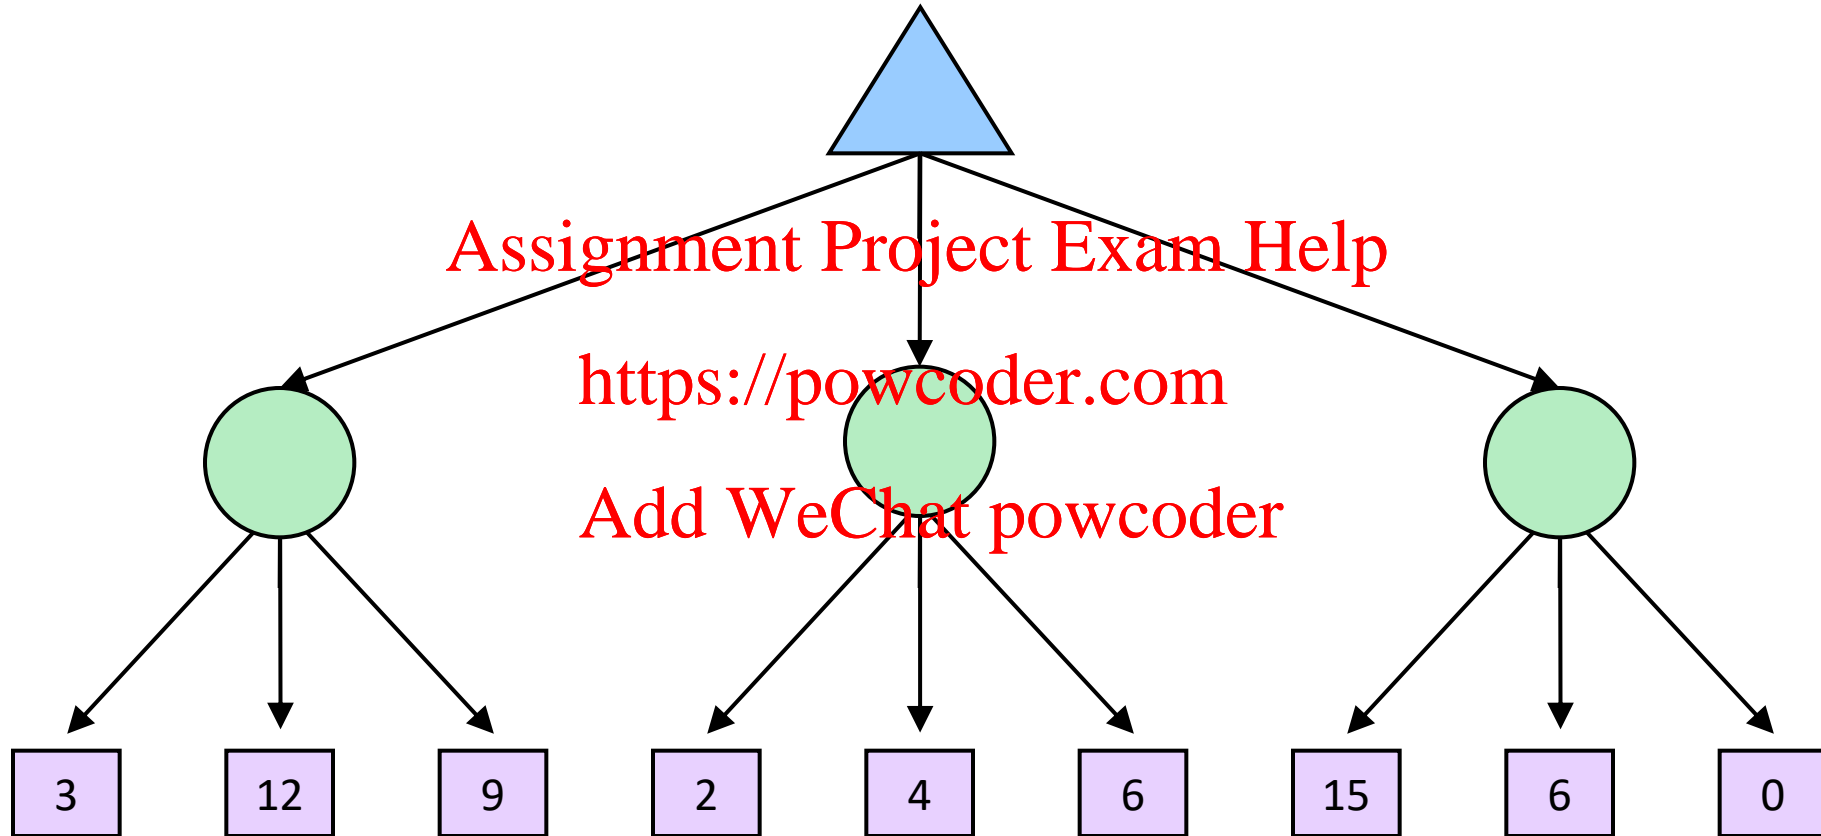
Add WeChat powcoder



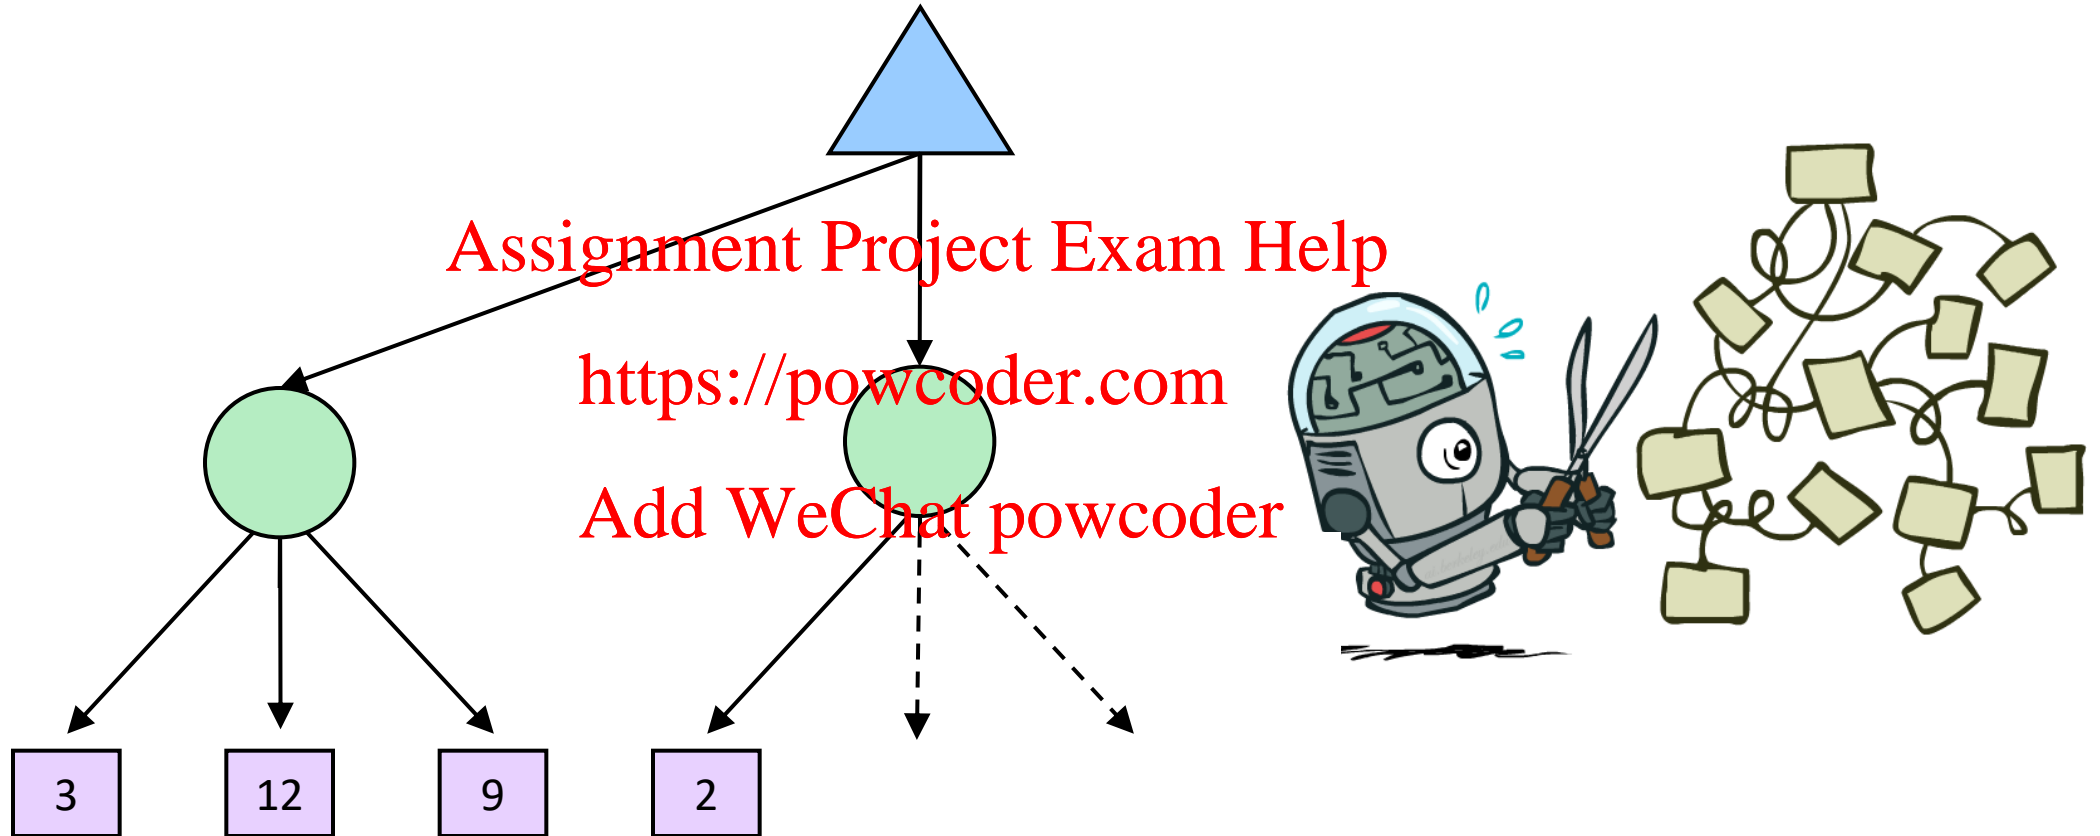
$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$



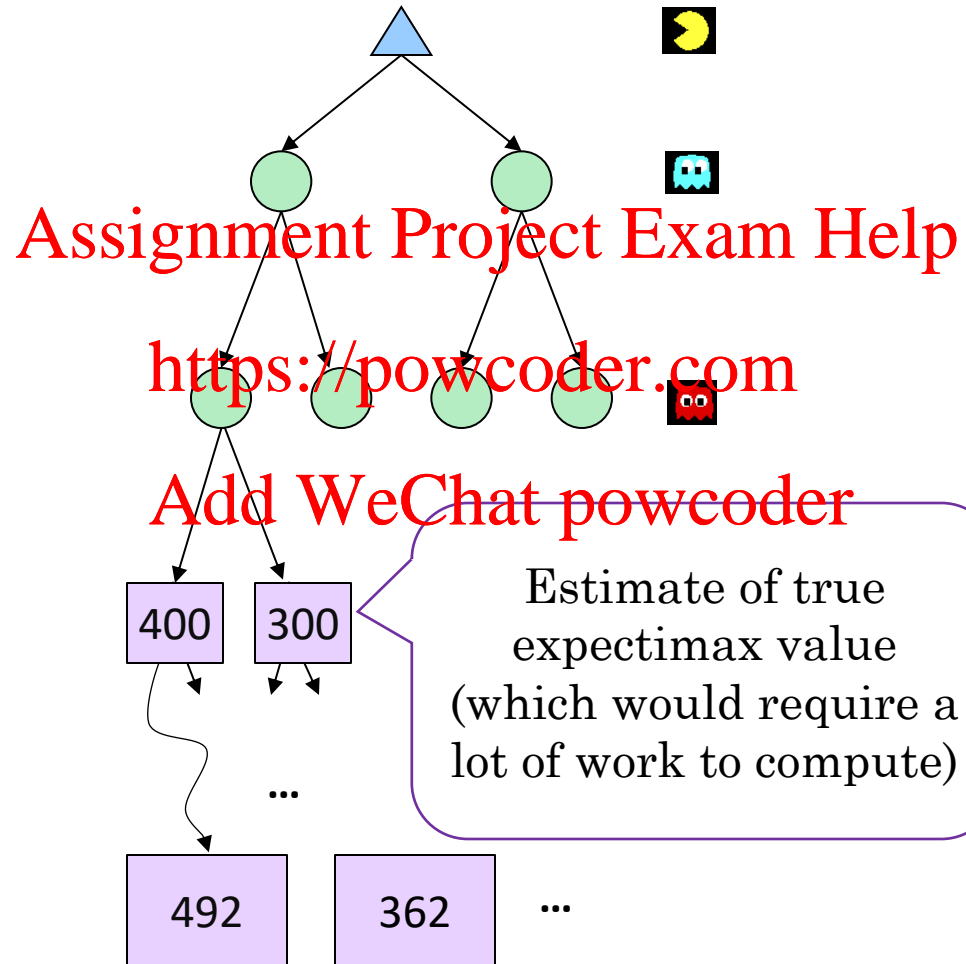
Expectimax Example



Expectimax Pruning?



Depth-Limited Expectimax




Probabilities

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Reminder: Probabilities

- A **random variable** represents an event whose outcome is unknown
- A **probability distribution** is an assignment of weights to outcomes

- Example: Traffic on freeway

- Random variable: T = whether there's traffic
- Outcomes: T in {none, light, heavy}
- Distribution: $P(T=\text{none}) = 0.25$, $P(T=\text{light}) = 0.50$, $P(T=\text{heavy}) = 0.25$

Add WeChat powcoder

- Some laws of probability (more later):
 - Probabilities are always non-negative
 - Probabilities over all possible outcomes sum to one
- As we get more evidence, probabilities may change:
 - $P(T=\text{heavy}) = 0.25$, $P(T=\text{heavy} \mid \text{Hour}=8\text{am}) = 0.60$
 - We'll talk about methods for reasoning and updating probabilities later



0.25



0.50



0.25



Reminder: Expectations

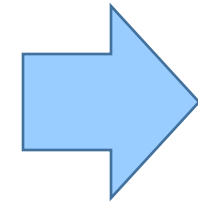
- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes

Assignment Project Exam Help

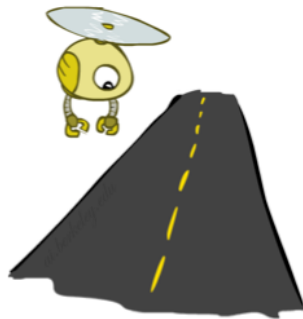

- Example: How long to get to the airport?

Time:	20 min	30 min	60 min
	x	x	x
Probability:	0.25	0.50	0.25

Add WeChat powcoder




35 min



What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state

- Model could be a simple uniform distribution (roll a die)
- Model could be sophisticated and require a great deal of computation
- We have a chance node for any outcome out of our control: opponent or environment
- The model might say that adversarial actions are likely!



- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes

Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!




Quiz: Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?

Assignment Project Exam Help

<https://powcoder.com>

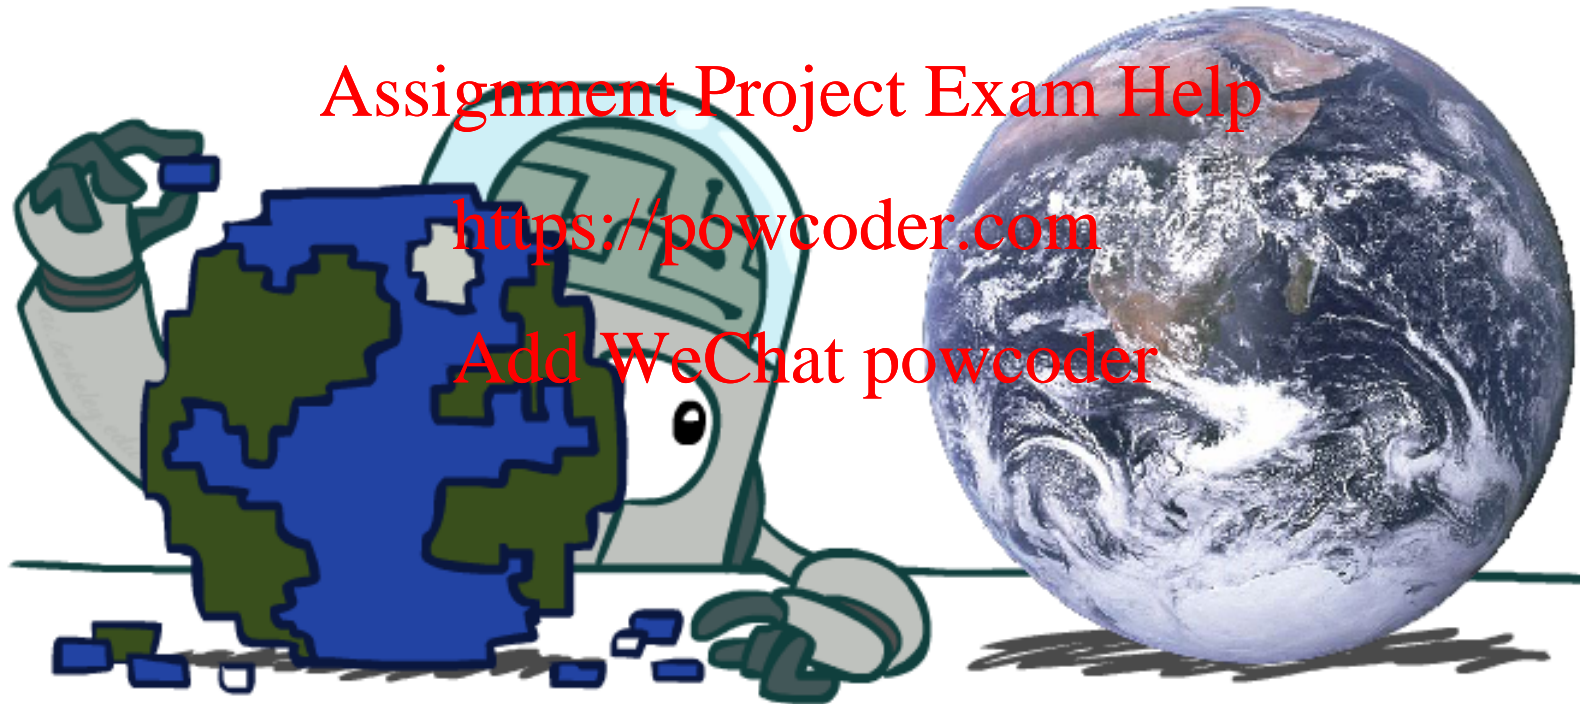
Add WeChat powcoder



- To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
- This kind of thing gets very slow very quickly
- Even worse if you have to simulate your opponent simulating you...
- ... except for minimax, which has the nice property that it all collapses into one game tree



Modeling Assumptions



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



The Dangers of Optimism and Pessimism

Dangerous Optimism

Assuming chance when the world is adversarial



Dangerous Pessimism

Assuming the worst case when it's not likely




Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assumptions vs. Reality



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


	Adversarial Ghost	Random Ghost
Minimax Pacman	1	1
Expectimax Pacman	1	1

Results from playing 5 games

Pacman used depth 4 search with an eval function that avoids trouble
Ghost used depth 2 search with an eval function that seeks Pacman



Assumptions vs. Reality



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

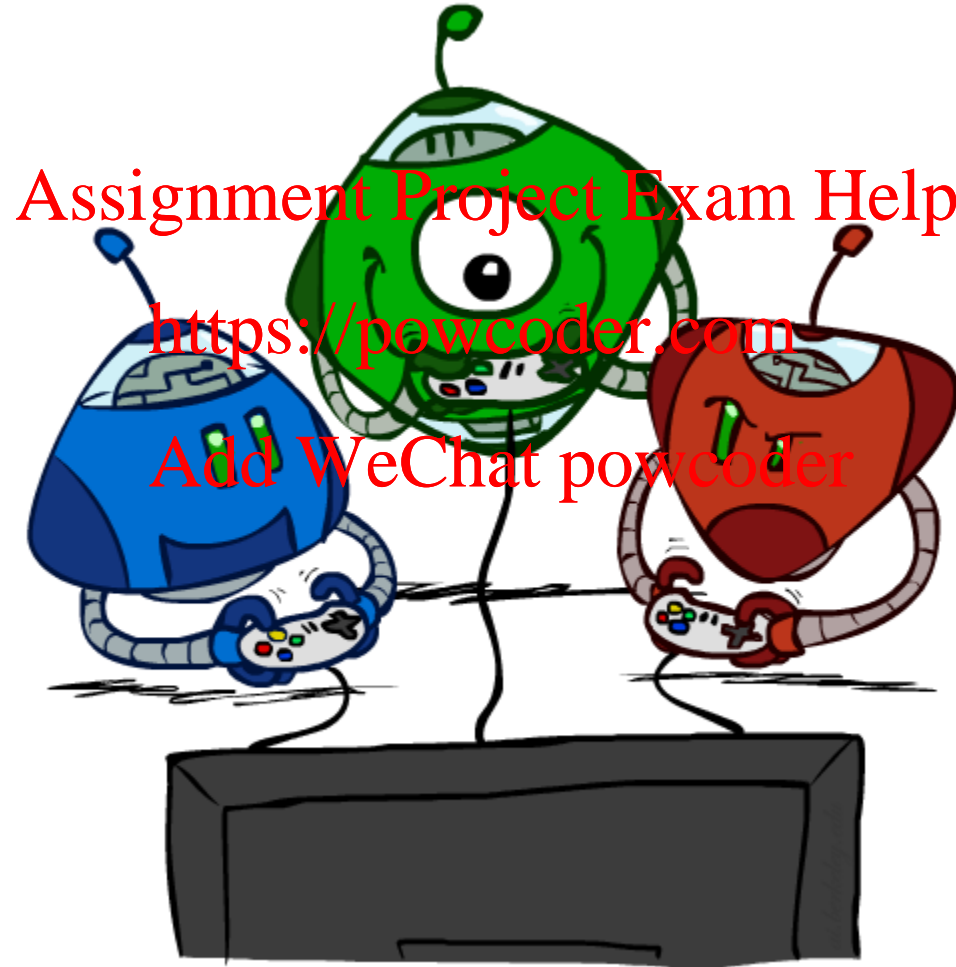
	Adversarial Ghost	Random Ghost
Minimax Pacman	Won 5/5 Avg. Score: 483	Won 5/5 Avg. Score: 493
Expectimax Pacman	Won 1/5 Avg. Score: -303	Won 5/5 Avg. Score: 503

Results from playing 5 games

Pacman used depth 4 search with an eval function that avoids trouble
Ghost used depth 2 search with an eval function that seeks Pacman




Other Game Types




Mixed Layer Types

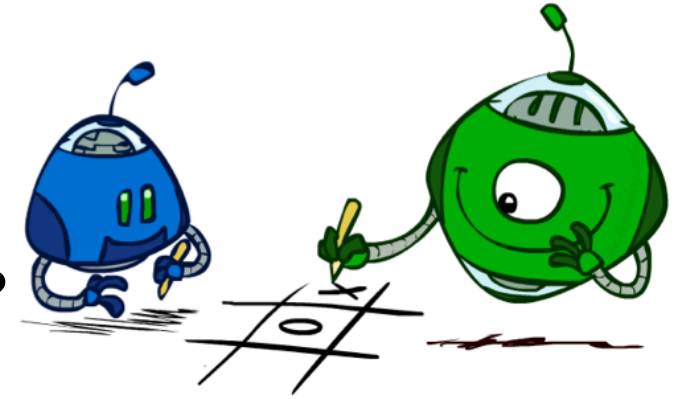
- E.g. Backgammon
- Expectiminimax
 - Environment is an extra “random agent” player that moves after each min/max agent
 - Each node computes the appropriate combination of its children




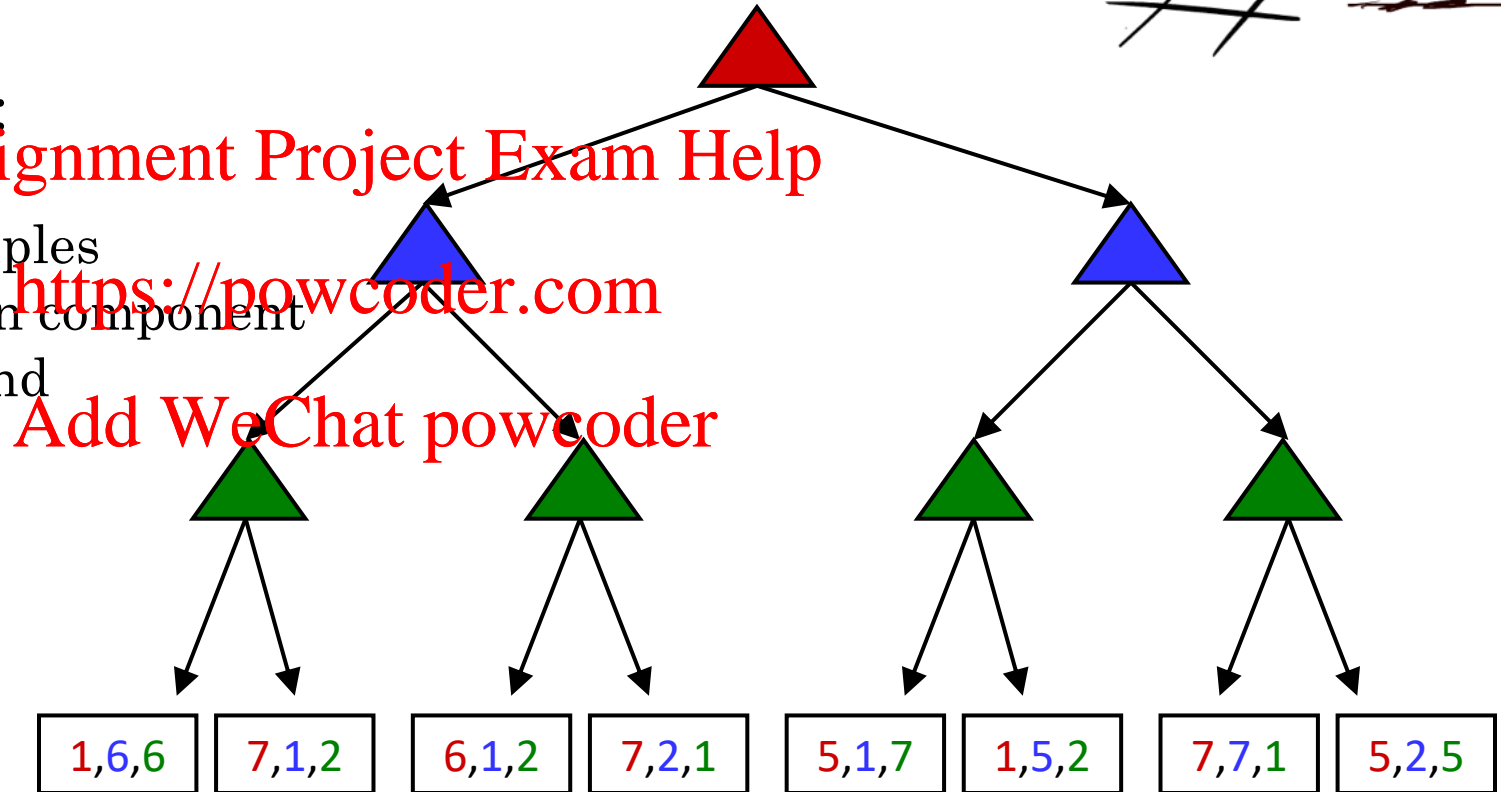
Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder




Multi-Agent Utilities



- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have utility tuples
 - Node values are also utility tuples
 - Each player maximizes its own component
 - Can give rise to cooperation and competition dynamically...



Utilities



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Maximum Expected Utility

- Why should we average utilities? Why not minimax?


- Principle of maximum expected utility:
 - A rational agent should choose the action that maximizes its expected utility, given its knowledge

Add WeChat powcoder

- Questions:
 - Where do utilities come from?
 - How do we know such utilities even exist?
 - How do we know that averaging even makes sense?
 - What if our behavior (preferences) can't be described by utilities?



What Utilities to Use?



Add WeChat powcoder

- For worst-case minimax reasoning, terminal function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - We call this **insensitivity to monotonic transformations**
- For average-case expectimax reasoning, we need *magnitudes* to be meaningful



Utilities

- Utilities are functions from outcomes (states of the world) to real numbers that describe an agent's preferences

Assignment Project Exam Help

<https://powcoder.com>


- Where do utilities come from?
 - In a game, may be simple (+1/-1)
 - Utilities summarize the agent's goals
 - Theorem: any “rational” preferences can be summarized as a utility function



Add WeChat powcoder



Utilities: Uncertain Outcomes



Preferences

- An agent must have preferences among:

- Prizes: A , B , etc.
- Lotteries: situations with uncertain prizes

$$L = [p, A; (1 - p), B]$$

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder


- Notation:

- Preference: $A \succ B$
- Indifference: $A \sim B$

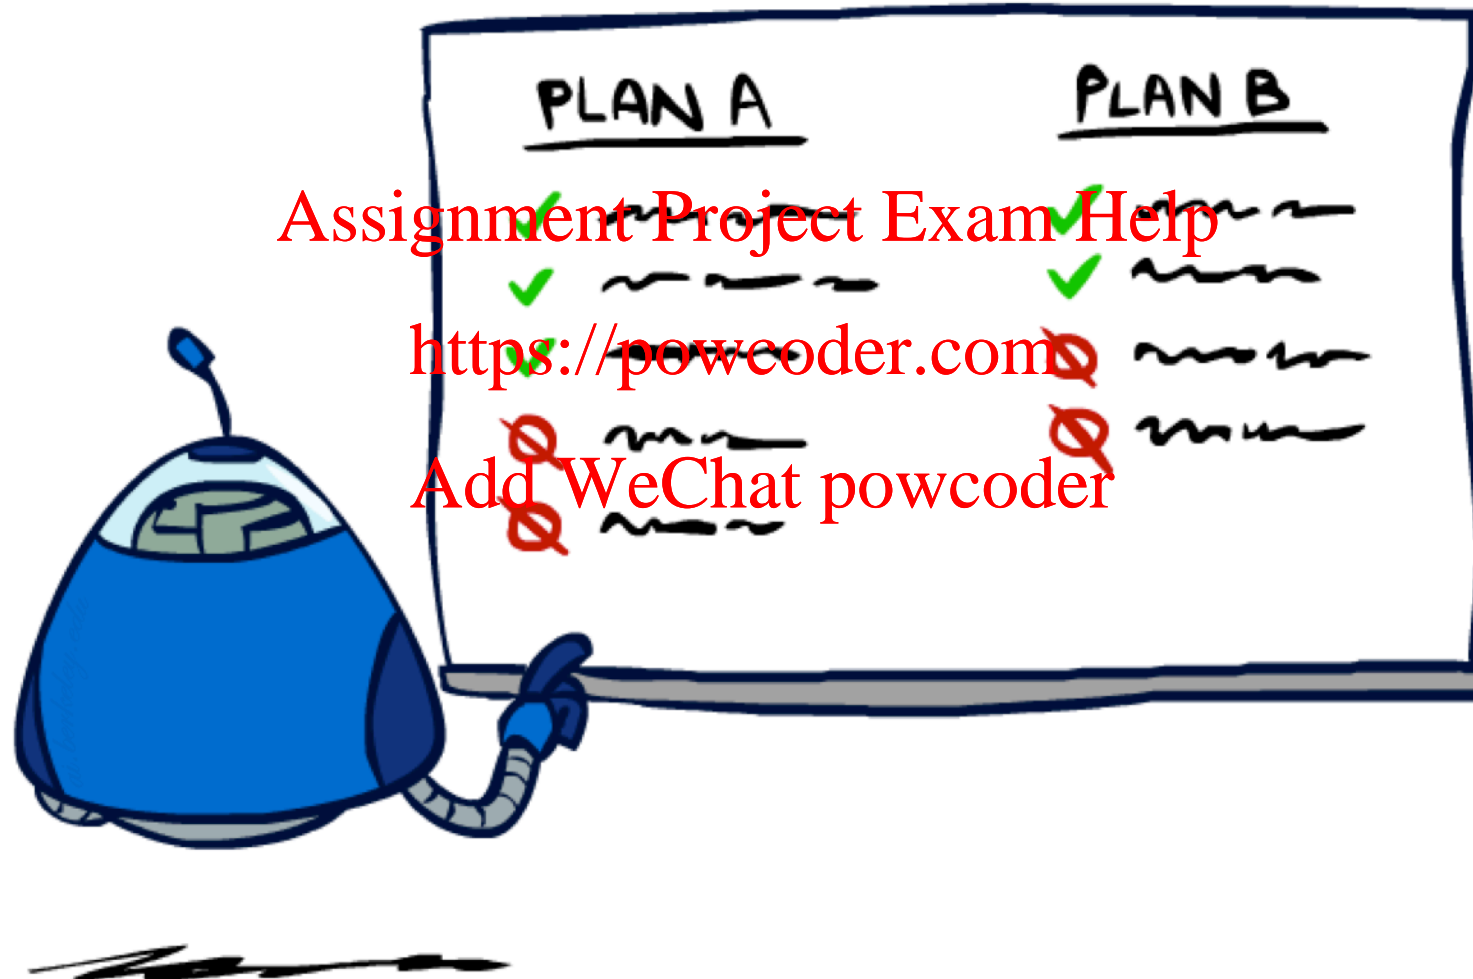
A Prize



A Lottery



Rationality



Rational Preferences


- We want some constraints on preferences before we call them rational, such as:

Axiom of Transitivity: $(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$

<https://powcoder.com>

Add WeChat powcoder

- For example: an agent with **intransitive preferences** can be induced to give away all of its money
 - If $B \succ C$, then an agent with C would pay (say) 1 cent to get B
 - If $A \succ B$, then an agent with B would pay (say) 1 cent to get A
 - If $C \succ A$, then an agent with A would pay (say) 1 cent to get C



Rational Preferences

The Axioms of Rationality

Orderability

$$(A \succ B) \vee (B \succ A) \vee (A \sim B)$$

Transitivity

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$

Continuity

$$A \succ B \succ C \Rightarrow \exists p [p, A; 1 - p, C] \sim B$$

Substitutability

$$A \sim B \Rightarrow [p, A; 1 - p, C] \sim [p, B; 1 - p, C]$$

Monotonicity

$$A \succ B \Rightarrow (p \geq q \Leftrightarrow [p, A; 1 - p, B] \succeq [q, A; 1 - q, B])$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Theorem: Rational preferences imply behavior describable as maximization of expected utility



MEU Principle

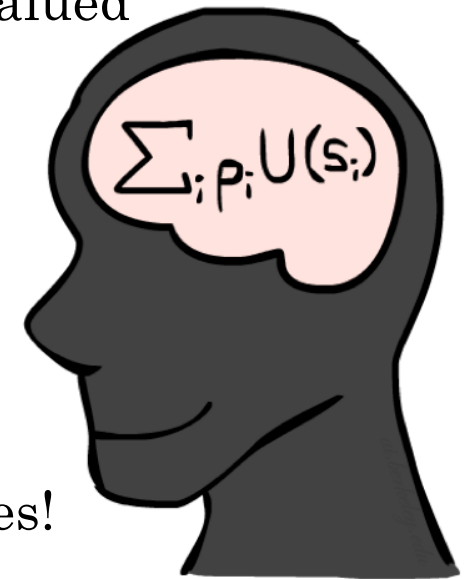
- Theorem [Ramsey, 1931; von Neumann & Morgenstern, 1944]
 - Given any preferences satisfying these constraints, there exists a real-valued function U such that:

$$U(A) \geq U(B) \Leftrightarrow A \succeq B$$


$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$$

Add WeChat powcoder

- I.e. values assigned by U preserve preferences of both prizes and lotteries!
- Maximum expected utility (MEU) principle:
 - Choose the action that maximizes expected utility



Human Utilities



Human Utilities

- Utilities map states to real numbers. Which numbers?
- Standard approach to assessment (elicitation) of human utilities:
 - Compare a prize A to a standard lottery L_p between
 - “best possible prize” u_+ with probability p
 - “worst possible catastrophe” u_- with probability $1-p$
 - Adjust lottery probability p until indifference: $A \sim L_p$
 - Resulting p is a utility in $[0,1]$



Human Utilities: Example

- A person is given the choice between 2 scenarios:
 - Guaranteed scenario: the person receives \$50
 - Uncertain scenario: a coin is flipped to decide the person receive \$100 or not.
- Which choice would that person make?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Risk Aversion

- **Risk averse**: would accept the guaranteed payment of (less than) \$50 rather than take the gamble

Assignment Project Exam Help

- **Risk neutral**: indifferent between the bet and the guaranteed \$50 payment

<https://powcoder.com>

Add WeChat powcoder

- **Risk seeking**: would accept the bet even when the guaranteed payment is more than \$50

Prospect Theory: Utility Function

- **Risk aversion:** convexity

- Risk averse regarding gain
- Risk seeking regarding loss

- **Loss aversion**

- Losses are felt more strongly than gains

- **Endowment effect**

- We value things we own more highly
- Reference point: differentiate gains and loss

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

